

Modern Elliptic Curve Cryptography

Benjamin Smith

SAC 2020 Summer School // October 19, 2020

Département d'Informatique, École polytechnique

The need for key agreement

Key agreement is a fundamental operation in cryptography.

It allows two principals to establish a shared secret key without prior contact.

The classic protocol for achieving this is **Diffie–Hellman Key Exchange**.

Historically, one of the first asymmetric crypto algorithms:

- **Public discovery:** Diffie and Hellman, 1976
- **Secret discovery:** GCHQ, UK, early 1970s.

For more generality and flexibility, we use **Key Encapsulation Mechanisms (KEMs)**; though sometimes we really do need Diffie–Hellman as a component in more complicated cryptosystems.

Keypairs

Keys come in matching (Public, Private) pairs.

Keypairs may be long-term (**static**) or single-use (**ephemeral**).

Every **public key** poses an individual mathematical **problem**;
the matching **private key** gives the **solution**.

In our discrete-log setting, keypairs present **DLP instances**
in a prime-order group $\mathcal{G} = \langle P \rangle$:

$$(\text{Public}, \text{Private}) = (Q, x) \quad \text{where} \quad Q = [x]P.$$


Diffie–Hellman key exchange

Goal: Alice and Bob want to establish a **shared secret** with no prior contact.

In Diffie–Hellman key exchange, we **combine secret scalars** from both sides using **composition** of scalar multiplications, which becomes **multiplication** of scalars.

Warning: *Diffie–Hellman has no built-in authentication!*

Diffie-Hellman key exchange (≤ 1976)



Correctness: $[a][b] = [b][a] = [ab]$ for all $a, b \in \mathbb{Z}$.

Alice & Bob now use a **KDF** (Key Derivation Function, e.g. HKDF) to derive a shared cryptographic key from the shared secret S .

The Diffie–Hellman problem


Diffie–Hellman security depends not (directly) on the DLP, but rather on the **Computational Diffie–Hellman Problem**:

Given $(P, Q_A = [x_A]P, Q_B = [x_B]P)$, compute $S = [x_A x_B]P$.

Relating the DLP and CDHP:

- Clearly the CDHP reduces immediately to the DLP:
you can solve a CDHP by solving a single DLP.
- There is a **conditional polynomial-time reduction** from the DLP to the CDHP (not obvious! Maurer–Wolf, ...)
- For the \mathcal{G} we use in practice, there is an **unconditional subexponential time reduction** from the DLP to the CDHP (Muzerou–Smart–Vercauteren).

Modern Diffie–Hellman key exchange



DH never directly uses the group structure on \mathcal{G} . All we need is a set \mathcal{G} and big randomly samplable sets A, B of efficiently computable functions $\mathcal{G} \rightarrow \mathcal{G}$ s.t.

- $[a][b] = [b][a]$ for all $[a] \in A$ and $[b] \in B$, and
- the corresponding CDHP is believed hard.

Diffie–Hellman does not need a group law, just scalar multiplication; so we can “drop signs” and work modulo \ominus .

Elliptic curves: work on x-line $\mathbb{P}^1 = \mathcal{E}/\langle \pm 1 \rangle$.

- The equivalence class $\{P = (x_P, y_P), \ominus P = (x_P, -y_P)\}$ is represented by the x-coordinate $\mathbf{x}(P) = x_P$.
- Projectively: $\mathbf{x}((X : Y : Z)) = (X : Z) \in \mathbb{P}^1$ when $Z \neq 0$, and $\mathbf{x}(\mathcal{O}) = \mathbf{x}((0 : 1 : 0)) = (1 : 0)$.

Advantage: save time and space by ignoring y .

Diffie–Hellman modulo signs

Diffie–Hellman does not need a group law, just scalar multiplication; so we can “drop signs” and work modulo \ominus .

The protocol is now

Alice computes $(a, x(P)) \mapsto x([a]P)$;

Bob computes $(b, x(P)) \mapsto x([b]P)$;

Alice computes $(a, x([b]P)) \mapsto x([a][b]P)$;


Bob computes $(b, x([a]P)) \mapsto x([b][a]P)$.

Question: is this well-defined?

*Question: How can we compute pseudo-scalar multiplication $(m, x(P)) \mapsto x([m]P)$ efficiently, **without using** \oplus ?*

Key fact: $\{x(P), x(Q)\}$ determines $\{x(P \oplus Q), x(P \ominus Q)\}$.

$\{x(P), x(Q)\}$ determines $\{x(P \ominus Q), x(P \oplus Q)\}$



Pseudo-group operations

Any 3 of $\mathbf{x}(P)$, $\mathbf{x}(Q)$, $\mathbf{x}(P \ominus Q)$, and $\mathbf{x}(P \oplus Q)$ determines the 4th, so we can define

Pseudo-addition:

$$\mathbf{xADD} : (\mathbf{x}(P), \mathbf{x}(Q), \mathbf{x}(P \ominus Q)) \mapsto \mathbf{x}(P \oplus Q)$$

Pseudo-doubling:

$$\mathbf{xDBL} : \mathbf{x}(P) \mapsto \mathbf{x}([2]P)$$

Bonus: it is easier to identify, isolate, and avoid special cases for \mathbf{xADD} than for \oplus .

In the following, we fix a **Montgomery curve**

$$\mathcal{E} : BY^2Z = X(X^2 + AXZ + Z^2)$$

with $A \neq \pm 2$ and $B \neq 0$ in \mathbb{F}_p .

Observe: we can convert to and from a short Weierstrass model for \mathcal{E} via $(X : Y : Z) \mapsto ((BX - AZ/3) : BY : Z)$, so all the theory we have already described transfers to this curve.

Notation: given points P and Q in $\mathcal{E}(\mathbb{F}_p)$, we write

$$P = (X_P : Y_P : Z_P),$$

$$Q = (X_Q : Y_Q : Z_Q),$$

$$P \oplus Q = (X_{\oplus} : Y_{\oplus} : Z_{\oplus}),$$

$$P \ominus Q = (X_{\ominus} : Y_{\ominus} : Z_{\ominus}).$$

Pseudo-addition on $\mathcal{E} : BY^2Z = X(X^2 + AXZ + Z^2)$:

$$\text{xADD} : (\mathbf{x}(P), \mathbf{x}(Q), \mathbf{x}(P \ominus Q)) \mapsto \mathbf{x}(P \oplus Q)$$

We use

$$(X_{\oplus} : Z_{\oplus}) = (Z_{\ominus} \cdot [U + V]^2 : X_{\ominus} \cdot [U - V]^2)$$

where

$$\begin{cases} U = (X_P - Z_P)(X_Q + Z_Q) \\ V = (X_P + Z_P)(X_Q - Z_Q) \end{cases}$$

Pseudo-doubling on \mathcal{E} : $BY^2Z = X(X^2 + AXZ + Z^2)$:

$$\text{xDBL} : \mathbf{x}(P) \longmapsto \mathbf{x}([2]P)$$

We use

$$(X_{[2]P} : Z_{[2]P}) = (Q \cdot R : S \cdot (R + \frac{A+2}{4}S))$$

where

$$\begin{cases} Q = (X_P + Z_P)^2, \\ R = (X_P - Z_P)^2, \\ S = 4X_P \cdot Z_P = Q - R. \end{cases}$$

We evaluate $x(P) \mapsto x([m]P)$ by combining **xADDs** and **xDBLs** using **differential addition chains**.

This means that **every** \oplus has summands with **known difference**.

Classic example: the **Montgomery ladder**.

The Montgomery ladder in a group

Algorithm 1: The Montgomery ladder in a group

Input: $m = \sum_{i=0}^{\beta-1} m_i 2^i$ and P

Output: $[m]P$

```
1  $(R_0, R_1) \leftarrow (0, P)$  // Invariant:  $R_1 = R_0 \oplus P$ 
2 for  $i$  in  $(\beta - 1, \dots, 0)$  do // invariant:  $R_0 = \lfloor m/2^i \rfloor P$ 
3   if  $m_i = 0$  then
4      $(R_0, R_1) \leftarrow ([2]R_0, R_0 \oplus R_1)$ 
5   else
6      $(R_0, R_1) \leftarrow (R_0 \oplus R_1, [2]R_1)$ 
7 return  $R_0$  //  $R_0 = [m]P, R_1 = [m + 1]P$ 
```

For each addition $R_0 \oplus R_1$, the difference $R_0 \ominus R_1$ is fixed (& known in advance!)
 \implies easy adaptation from \mathcal{E} to \mathbb{P}^1 .

The Montgomery ladder with pseudo-operations

Algorithm 2: The Montgomery ladder on the x-line \mathbb{P}^1

Input: $m = \sum_{i=0}^{\beta-1} m_i 2^i$ and $\mathbf{x}(P)$

Output: $\mathbf{x}([m]P)$

```
1  $(x_0, x_1) \leftarrow (\mathbf{x}(0), \mathbf{x}(P))$ 
2 for  $i$  in  $(\beta - 1, \dots, 0)$  do
3   if  $m_i = 0$  then
4      $(x_0, x_1) \leftarrow (\mathbf{xDBL}(x_0), \mathbf{xADD}(x_0, x_1, \mathbf{x}(P)))$ 
5   else
6      $(x_0, x_1) \leftarrow (\mathbf{xADD}(x_0, x_1, \mathbf{x}(P)), \mathbf{xDBL}(x_1))$ 
7 return  $x_0$                                      //  $x_0 = \mathbf{x}([m]P), R_1 = \mathbf{x}([m+1]P)$ 
```

The loop invariant is $(x_0, x_1) = (\mathbf{x}(\lfloor m/2^i \rfloor P), \mathbf{x}(\lfloor m/2^i \rfloor + 1)P)$.

X25519 is the state-of-the-art Diffie–Hellman key-exchange algorithm, standardized for TLS 1.3 (and other applications).

- Based on Bernstein's **Curve25519** software (2006)
- Formalized in **RFC7748**, *Elliptic curves for security* (2016)

It is a massive upgrade on traditional ECDH (used e.g. in TLS ≤ 1.2), which was based on NIST's standard prime-order curves.

Bernstein (PKC 2006) defined an elliptic curve

$$\mathcal{E} : Y^2Z = X(X^2 + 486662 \cdot XZ + Z^2) \quad \text{over } \mathbb{F}_p$$

where $p = 2^{255} - 19$. This curve is known as **Curve25519**.

The curve has order $\#\mathcal{E}(\mathbb{F}_p) = 8r$, where r is prime.

Bernstein (PKC 2006) defined an elliptic curve

$$\mathcal{E} : Y^2Z = X(X^2 + 486662 \cdot XZ + Z^2) \quad \text{over } \mathbb{F}_p$$

where $p = 2^{255} - 19$. This curve is known as **Curve25519**.

The curve has order $\#\mathcal{E}(\mathbb{F}_p) = 8r$, where r is prime.

If we let B be any nonsquare in \mathbb{F}_p , then the **quadratic twist**

$$\mathcal{E}' : B \cdot Y^2Z = X(X^2 + 486662 \cdot XZ + Z^2) \quad \text{over } \mathbb{F}_p$$

has order $\#\mathcal{E}'(\mathbb{F}_p) = 4r'$, where r' is prime.

The X25519 function

The **X25519 function** maps $\mathbb{Z}_{\geq 0} \times \mathbb{F}_p$ into \mathbb{F}_p , via

$$(m, u) \mapsto u_m := x_m \cdot z_m^{(p-2)}$$

where $(x_m : * : z_m) = [m](u : * : 1) \in \mathcal{E}(\mathbb{F}_p) \cup \mathcal{E}'(\mathbb{F}_p)$.

Note: generally $z_m \neq 0$, in which case $(u_m : * : 1) = [m](u : * : 1)$ in $\mathcal{E}(\mathbb{F}_p)$ or $\mathcal{E}'(\mathbb{F}_p)$.

Exercise: show that for any given u , inverting $(m, u) \mapsto u_m$ amounts to solving a discrete logarithm in either $\mathcal{E}(\mathbb{F}_p)$ or $\mathcal{E}'(\mathbb{F}_p)$.

Diffie–Hellman with X25519

The global public “base point” is $u_1 = 9 \in \mathbb{F}_p$.

The point $(u_1 : * : 1)$ has order r in $\mathcal{E}(\mathbb{F}_p)$ (remember: r is a 252-bit prime).

The “scalars” are integers in $S = \{2^{254} + 8i : 0 \leq i < 2^{251}\}$.

Key generation:

Alice samples secret $a \in S$, computes $A := u_a = \text{X25519}(a, u_1)$, publishes A .

Bob samples secret $b \in S$, computes $B := u_b = \text{X25519}(b, u_1)$, publishes B .

Shared secret:

Alice computes $S = u_{ab} = \text{X25519}(a, B)$

Bob computes $S = u_{ba} = \text{X25519}(b, A)$

Side-channel concerns

Our algorithms must anticipate basic **side-channel attacks** (*especially timing attacks and power analysis*).

Diffie–Hellman implementations must be “**uniform**” and “**constant-time**” with respect to the secret scalars:

- No branching on bits of secrets
eg. No **if**(**m** == **0**): ... with m_i secret
- No memory accesses indexed by (bits of) secrets
(eg. No **x** = **T**[**m**] where m is secret)

What we want is to have *exactly the same sequence of computer instructions* for every possible secret input.

Towards a uniform/constant-time Montgomery ladder

Algorithm 3: The Montgomery ladder for X25519

Input: $m = \sum_{i=0}^{\beta-1} m_i 2^i$ and $x = \mathbf{x}(P)$ with P in $\mathcal{E}(\mathbb{F}_p)$ or $\mathcal{E}'(\mathbb{F}_p)$

Output: $\mathbf{x}([m]P)$

```
1  $\mathbf{u} \leftarrow (x, 1)$ 
2  $(\mathbf{x}_0, \mathbf{x}_1) \leftarrow ((1, 0), \mathbf{u})$ 
3 for  $i$  in  $(\beta - 1, \dots, 0)$  do
4   if  $m_i = 0$  then           // Remove branch using conditional swaps
5      $(\mathbf{x}_0, \mathbf{x}_1) \leftarrow (\text{xDBL}(\mathbf{x}_0), \text{xADD}(\mathbf{x}_0, \mathbf{x}_1, \mathbf{u}))$  // Uniform xDBL, xADD
6   else
7      $(\mathbf{x}_0, \mathbf{x}_1) \leftarrow (\text{xADD}(\mathbf{x}_0, \mathbf{x}_1, \mathbf{u}), \text{xDBL}(\mathbf{x}_1))$ 
8 return  $\mathbf{x}_0$ 
```

Conditional swap algorithms

Algorithm 4: Conditional swap: parallel bit operations

1 Function *SWAP*

Input: $b \in \{0, 1\}$ and (x_0, x_1)

Output: (x_0, x_1) if $b = 0$, (x_1, x_0) if $b = 1$

2 $v \leftarrow b \text{ and } (x_0 \text{ xor } x_1)$

3 return $(x_0 \text{ xor } v, x_1 \text{ xor } v)$

Algorithm 5: Conditional swap: arithmetic operations

1 Function *SWAP*

Input: $b \in \{0, 1\}$ and (x_0, x_1)

Output: (x_0, x_1) if $b = 0$, (x_1, x_0) if $b = 1$

2 return $((1 - b)x_0 + bx_1, bx_0 + (1 - b)x_1)$

Towards a uniform/constant-time Montgomery ladder

Algorithm 6: The Montgomery ladder for X25519

Input: $m = \sum_{i=0}^{\beta-1} m_i 2^i$ and $x = \mathbf{x}(P)$ with P in $\mathcal{E}(\mathbb{F}_p)$ or $\mathcal{E}'(\mathbb{F}_p)$

Output: $\mathbf{x}([m]P)$

```
1  $\mathbf{u} \leftarrow (x, 1)$ 
2  $(\mathbf{x}_0, \mathbf{x}_1) \leftarrow ((1, 0), \mathbf{u})$ 
3 for  $i$  in  $(\beta - 1, \dots, 0)$  do
4    $(\mathbf{x}_0, \mathbf{x}_1) \leftarrow \text{SWAP}(m_i, (\mathbf{x}_0, \mathbf{x}_1))$ 
5    $(\mathbf{x}_0, \mathbf{x}_1) \leftarrow (\mathbf{xDBL}(\mathbf{x}_0), \mathbf{xADD}(\mathbf{x}_0, \mathbf{x}_1, \mathbf{u}))$ 
6    $(\mathbf{x}_0, \mathbf{x}_1) \leftarrow \text{SWAP}(m_i, (\mathbf{x}_0, \mathbf{x}_1))$ 
7 return  $\mathbf{x}_0$ 
```

Clearly we can halve the number of conditional swaps (try doing this!).

Elliptic curve signature schemes

Keys come in matching (Public,Private) pairs.

Every public key poses an individual mathematical problem; the matching private key gives the solution.

Here, keypairs present an instances of the DLP in $\mathcal{G} = \langle P \rangle$:

$$(\text{Public}, \text{Private}) = (Q, x) \quad \text{where} \quad Q = [x]P$$

where P is some fixed generator of \mathcal{G} .

Having an **identity** means being distinguishable, from everyone else.

***Question:** what is identity, in a cryptographic sense?*

Having an **identity** means being distinguishable, from everyone else.

Question: what is identity, in a cryptographic sense?

Identity means **holding the private key** corresponding to a bound public key.

We want to cryptographically **authenticate communicating parties** in a protocol.

That is: we want to know that we are communicating with someone holding the secret x corresponding to some public $Q = [x]P$.

Recall that in **symmetric** crypto, MACs and AEAD can authenticate **data**, but **not communicating parties**. The reason for this is simple: in symmetric crypto, *both sides hold the same secret*—and a shared identity is no identity.

How do you prove your identity?

In our setting, you assert or claim an identity by **binding** to (that is, publishing and committing to) a public key Q from a keypair $(Q = [x]P, x)$.

Prove your identity \iff prove you know x .


To formalize this, we introduce three characters:

Prover wants to *prove* their identity

Verifier wants to *verify* the identity of Prover

Simulator wants to impersonate Prover


Ineffective identification



1. Verifier challenges;
2. Prover returns x ;
3. Verifier accepts iff $[s]P = Q$.

Problem: Prover no longer has an identity, because they gave away their secret x .

Ineffective identification




1. Verifier challenges;
2. Prover returns x ;
3. Verifier accepts iff $[s]P = Q$.

Problem: Prover no longer has an identity, because they gave away their secret x .

Solution: hide long-term secrets with disposable one-shot (ephemeral) secrets.

Using ephemeral keys




1. Prover generates an *ephemeral* keypair (R, r) , commits to R ;
2. Verifier challenges;
3. Prover sends R and $s = x + r$ to Verifier.

Note: s reveals nothing about x , because r is random

Verifier accepts because $[s]P = [x]P + [r]P = Q + R$.

Using ephemeral keys

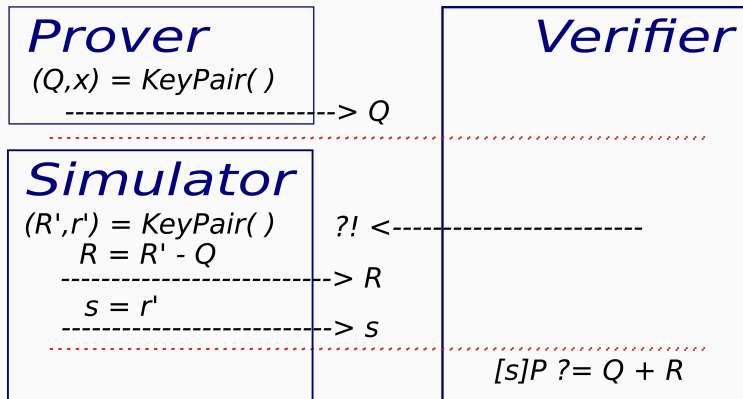


1. Prover generates an *ephemeral* keypair (R, r) , commits to R ;
2. Verifier challenges;
3. Prover sends R and $s = x + r$ to Verifier.

Note: s reveals nothing about x , because r is random

Verifier accepts because $[s]P = [x]P + [r]P = Q + R$. **Problem:** easy to cheat!

Cheating: Simulator can easily impersonate the Prover



Verifier accepts because $[s]P = [r']P = R' = Q + R$

Note: Simulator never knows x —nor the log of R , because then they would know x !

Detecting cheating

How can Verifier detect this cheating, and thus distinguish between Prover and Simulator?

- Prover**
- sends $s = x + r = \log(Q + R)$,
 - knows *both* $x = \log(Q)$ and $r = \log(R)$.

- Simulator**
- sends $s = \log(Q + R)$,
 - knows *neither* $x = \log(Q)$ *nor* $r = \log(R)$.

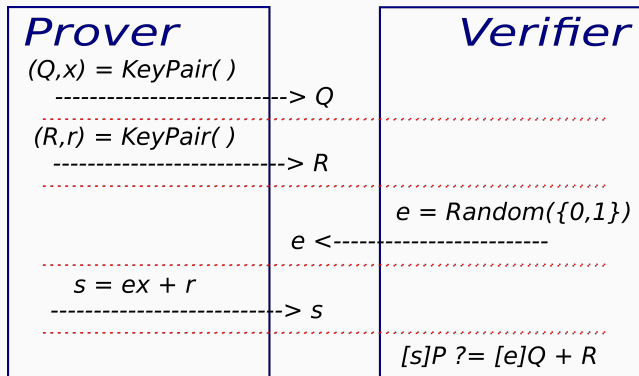
The difference: knowledge of x , and knowledge of r .

Verifier can't ask for x .

If she asks for the ephemeral secret $r = \log(R)$ while knowing s , then that would also reveal x .

Solution: let Verifier ask for **either** s **or** r ,
and check either $[s]P = Q + R$ or $[r]P = R$ accordingly.

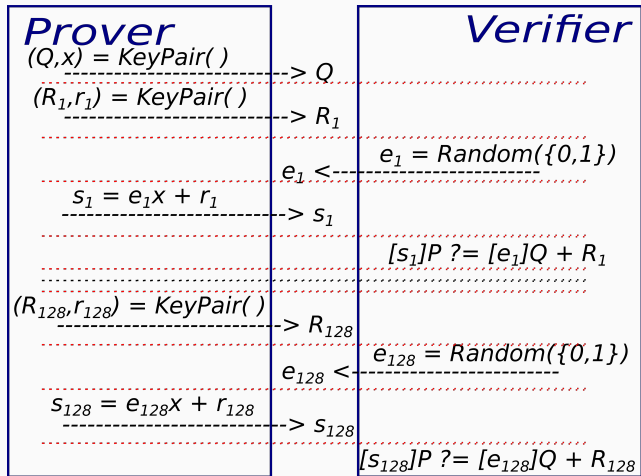
- correct $s \implies$ I know x , *if* I am honest
- correct $r \implies$ I was honest, but *not* that I know x



To cheat, Simulator must guess/anticipate e : 50% chance.

So repeat until Verifier is satisfied it's Prover (say 128 rounds).

128 rounds later...



The Chaum–Evertse–Graaf ID protocol is **zero knowledge**:

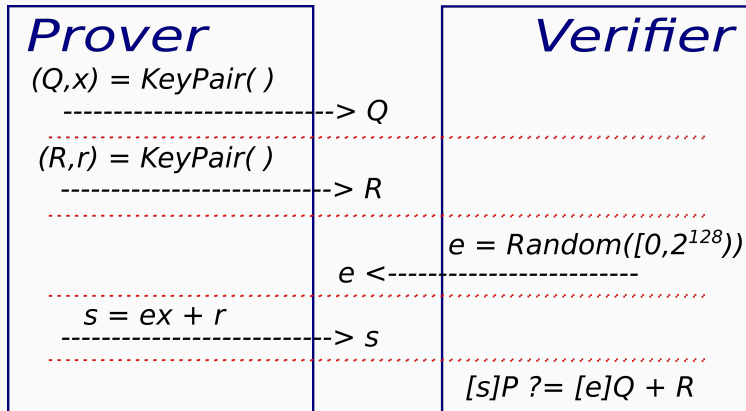
- Verifier is convinced that Prover holds the secret x , but
- Verifier learns absolutely nothing about the value of x .

This holds independent of the number of rounds carried out.

Running 128 rounds of Chaum–Evertse–Graaf ID is extremely inconvenient:

1. too many **interactive rounds** of communication
(128 challenges and responses),
2. too much **bandwidth**
(128×256 -bit group elements and 128×256 -bit scalars)
3. too much **computation** on each side
(128×256 -bit scalar multiplications for both parties!)

Schnorr identification (1991): “**parallelise**” the 128 rounds, replacing 128 one-bit challenges with one 128-bit challenge.



Note: s reveals nothing about x , because r is random.

One round. Scalar mults: Prover = 1×256 -bit, Verifier = $1 \times 256 + 1 \times 128$ -bit.

The Schnorr ID protocol is **not zero knowledge**.

The Schnorr ID protocol is **not zero knowledge**.

- The triple (Q, s, e) is a particular solution of the equation $R = [s]P \ominus [e]Q$,
- so Verifier can choose e as a function of Q .

Still, it is (reasonably) clear that Verifier does not learn anything **useful** about x from the protocol.

A **digital signature** is a **non-interactive proof** that the Signer witnessed (created, saw) some data.

Authenticity, message integrity, non-repudiability:

- only the Signer could have created it;
- the Signer could not have created it from any other data; and
- only the Signer's public key is needed to *verify* it.

The Fiat-Shamir transform

We build **Schnorr signatures** from the Schnorr ID scheme by applying the **Fiat-Shamir transform**:

1. make the ID scheme non-interactive, and
2. have the signer identify themselves *to the data* (!)

Formally: **Fiat-Shamir** transforms an interactive proof with public randomness into a non-interactive proof, by replacing the verifier with a cryptographic hash function applied to the protocol's transcript.

Step one: making Schnorr ID non-interactive

Intuition: the hash of R is unpredictable and random-looking, so it can stand in for a true random challenge.

Prover

$(Q, x) = \text{KeyPair}()$

-----> Q

$(R, r) = \text{KeyPair}()$

-----> R

$e = \text{Hash}(R)$

$s = ex + r$

-----> s

Verifier

$e = \text{Hash}(R)$

$[s]P \stackrel{?}{=} [e]Q + R$

Step two: “identify to the data”

Prover

$(Q, x) = \text{KeyPair}()$

-----> Q

$(R, r) = \text{KeyPair}()$

$e = \text{Hash}(R)$

-----> e

$s = ex + r$

-----> s

Verifier

$R = [s]P - [e]Q$

$e \stackrel{?}{=} \text{Hash}(R)$

Generally (especially if $\mathcal{G} = \mathbb{F}^\times$) the hash e is smaller than R , so we can send (e, s) instead of (R, s) .

Schnorr signatures (1991)

Signer

$(Q, x) = \text{KeyPair}()$

----- $\rightarrow Q$

$(R, r) = \text{KeyPair}()$

$e = \text{Hash}(R, M)$

$M \leftarrow \text{Message}$

----- $\rightarrow e$

$s = ex + r$

----- $\rightarrow s$

Verifier

$R = [s]P - [e]Q$

$e \stackrel{?}{=} \text{Hash}(R, M)$

The hash function must provide 128 bits of **prefix-second-preimage** resistance.

Traditionally **no need for collision resistance**, though you might want it to protect against attacks on multiple keys.

Schnorr signatures

Schnorr signatures are proven secure in the **random oracle model**, though not in the **standard model**.

Schnorr **patented** his signature scheme; as a result, few people actually used it. *(Instead, we ended up with the inferior DSA and ECDSA protocols...)*

The patent **expired in 2008**.

This has led to the rise of contemporary Schnorr variants like **EdDSA**.

Bonus Track 1:
EdDSA signatures

EdDSA (Bernstein–Duif–Lange–Schwabe–Yang, 2012)

Fix a 2β -bit hash function H and a secure elliptic curve \mathcal{E}/\mathbb{F}_p such that $\mathcal{E}(\mathbb{F}_p)$ contains a $\approx \beta$ -bit prime-order subgroup $\mathcal{G} = \langle P \rangle$.

Key generation: choose a random β -bit string, k .

Let x and y be the β -bit strings s.t. $x \parallel y = H(k)$.

Public key: $Q = [x]P$. **Secret key:** k (not x).

Signing a message M under x :

let $r = H(y \parallel M)$, $R = [r]P$, $s = r + H(R \parallel Q \parallel M)x$.

Signature: (R, s) .

Verifying a putative signature (R, s) on M under Q :

accept iff $R = [s]P \ominus [H(R \parallel Q \parallel M)]Q$.

Question: what are the important differences between EdDSA and plain Schnorr?

Algorithm 7: EdDSA Key Generation

Output: $(Q, k) \in \mathcal{E}(\mathbb{F}_p) \times \{0, 1\}^\beta$

- 1 $k \leftarrow \text{Random}(\{0, 1\}^\beta)$
 - 2 $x \parallel y \leftarrow H(k)$ // Both x and y are in $\{0, 1\}^\beta$
 - 3 $Q \leftarrow [x]P$
 - 4 **return** (Q, k)
-

Algorithm 8: EdDSA signing

Input: Message $m \in \{0, 1\}^*$, secret key k

Output: Signature $\sigma = (R, s) \in \mathcal{G} \times \mathbb{Z}/N\mathbb{Z}$

- 1 $x \parallel y \leftarrow H(k)$
 - 2 $r \leftarrow H(y \parallel M)$
 - 3 $R \leftarrow [r]P$
 - 4 $s \leftarrow r + H(R \parallel Q \parallel M)x$
 - 5 **return** $\sigma = (R, s)$
-

Algorithm 9: EdDSA verification

Input: Message $m \in \{0, 1\}^*$, signature $\sigma = (R, s) \in \mathcal{G} \times \mathbb{Z}/N\mathbb{Z}$,
public key $Q = [x]P$

Output: True or False

```
1  $h \leftarrow H(R \parallel Q \parallel M)$ 
2  $R' = [s]P \oplus [-h]Q$  // Multiexponentiation
3 if  $R' = R$  then
4   | return True
5 else
6   | return False
```

Bonus Track 2:
Multiscalar multiplication

Multiscalar multiplication

For k bits of security ($2k$ -bit curve), intensive computations:

Signing 1x scalar multiplication by a $2k$ -bit scalar

Schnorr Verif. 1x scalar mult. by a k -bit scalar,
plus 1x scalar mult. by a $2k$ -bit scalar

EdDSA Verif. 2x scalar mults by $2k$ -bit scalars

The **cost of verification** is the chief **drawback** of EC signatures.

Multiexponentiation: a standard trick to **merge** “parallel” scalar multiplications into one, reducing verification cost...

Exercise: *compare the expected cost of key generation, signing, and verification for RSA and elliptic-curve signatures.*

Multisexponentiation with Straus' algorithm

Algorithm 10: Generic binary multiexponentiation

1 **Function** *Multiexponentiation*

Input: $(a = \sum_{i=0}^{\beta-1} a_i 2^i, b = \sum_{i=0}^{\beta-1} b_i 2^i)$ in $[0..2^\beta)^2$, P, Q in \mathcal{G}

Output: $[a]P \oplus [b]Q$

2 $(T_{0,0}, T_{1,0}, T_{0,1}, T_{1,1}) \leftarrow (0_{\mathcal{G}}, P, Q, P \oplus Q)$

3 $R \leftarrow 0_{\mathcal{G}}$

4 **for** $i = \beta - 1$ **down to** 0 **do**

5 $R \leftarrow [2]R \oplus T_{a_i, b_i}$

6 **return** R

There exist **uniform** and also **differential** (x-only) variants.

#doubles same as plain exponentiation by $\max(|a|, |b|)$. What about #adds?

Bonus Track 3:
ECDSA

ECDSA is the standard Elliptic Curve Digital Signature Algorithm.

Targeting k bits of security:

Fix a $2k$ -bit hash function H and a secure elliptic curve \mathcal{E}/\mathbb{F}_p such that $\mathcal{E}(\mathbb{F}_p)$ contains a $\approx 2k$ -bit prime-order subgroup $\mathcal{G} = \langle P \rangle \cong \mathbb{Z}/N\mathbb{Z}$.

Public-private keypairs are standard DLP instances ($Q = [x]P, x$).

Algorithm 11: ECDSA signing

Input: Message $m \in \{0, 1\}^*$, secret key x

Output: Signature $\sigma = (s, t) \in [1, N - 1]^2$

```
1  $z \leftarrow H(m)$ 
2  $(R = [r]P, r) \leftarrow \text{KeyGen}()$  // Ephemeral keypair
3  $s \leftarrow x(R) \bmod N$  // Viewing x-coord  $x(R) \in \mathbb{F}_p$  as an integer!
4 if  $s = 0$  then go to Line 2
5  $t \leftarrow r^{-1}(z + sx) \pmod{N}$ 
6 if  $t = 0$  then go to Line 2
7 return  $\sigma = (s, t)$ 
```

Algorithm 12: ECDSA signing

Input: Message $m \in \{0, 1\}^*$, signature $\sigma = (s, t) \in [1, N - 1]^2$, public key $Q = [x]P$

Output: True or False

```
1  $z \leftarrow H(m)$ 
2  $u \leftarrow zt^{-1} \pmod{N}$ 
3  $v \leftarrow st^{-1} \pmod{N}$ 
4  $S \leftarrow [u]P \oplus [v]Q$                                 // Multiexponentiation
5 if  $s \equiv x(S) \pmod{N}$  then                             // Viewing  $x(S)$  as an integer!
6   | return True
7 else
8   | return False
```

Disadvantages of ECDSA

ECDSA is **harder to implement correctly** than plain Schnorr or EdDSA.

It is also **harder to prove** things about ECDSA than it is for Schnorr or EdDSA, because it involves weird operations like changing moduli (from p to N , etc.)

For ECDSA, Schnorr, and EdDSA, the keys and **signatures are exceptionally small** (this is the USP of EC signatures).

Main drawback: **verification is relatively slow** (even with multiexponentiation).